

#1. Load both CSV files into separate Pandas DataFrames.

```
import pandas as pd
import numpy as np
```

```
customers_df = pd.read_csv("customers_200_rows.csv")
sales_df = pd.read_csv("sales_200_rows.csv")
```

#2. Display the first 5 and last 5 rows of each DataFrame.

```
print(customers_df.head())
print(customers_df.tail())
print(sales_df.head())
print(sales_df.tail())
```

	customer_id	name	email
0	1001	Norma Fisher	ysullivan@yahoo.com
1	1002	Susan Wagner	katelynmontgomery@yahoo.com
2	1003	Dr. Stephanie Collins	thomas15@stewart-bowman.com
3	1004	Joseph Brown	cortezraymond@garrett.com
4	1005	Amy Stark	lindathomas@west.net

	country	signup_date
0	Lesotho	2023-12-20
1	United States of America	2024-09-16
2	Mexico	2024-06-22
3	Ecuador	2023-10-30
4	Venezuela	2024-07-11

	customer_id	name	email
195	1196	Robin Schroeder	roberjonjulie@phillips-daniel.biz
196	1197	Madison Hicks	williamsalexis@beasley.biz
197	1198	Emily Weiss	vschneider@williams.com
198	1199	Brandi Simon	isullivan@gmail.com
199	1200	Brianna Pugh	briannajackson@ray.com

	country	signup_date
195	Estonia	2023-08-14
196	Slovenia	2024-06-28
197	Australia	2024-12-28
198	Samoa	2024-05-04
199	El Salvador	2023-05-06

	order_id	customer_id	order_date	product	category
0	5001	1071	2023-09-19	Tablet	Electronics
4					
1	5002	1035	2022-10-01	Headphones	Accessories

```

1
2      5003      1093  2023-04-01      Webcam  Accessories
1
3      5004      1057  2023-07-12  Smartphone  Electronics
1
4      5005      1100  2023-03-13      Laptop  Electronics
2

```

```

    price_per_unit
0      399.00
1      89.99
2      59.00
3     599.00
4     789.99
    order_id  customer_id  order_date  product  category
quantity \
195      5196      1011  2022-05-06   Printer  Electronics
3
196      5197      1045  2022-12-11  Keyboard  Accessories
1
197      5198      1052  2022-12-05   Laptop  Electronics
4
198      5199      1051  2023-08-02    Mouse  Accessories
2
199      5200      1008  2023-01-05   Charger  Accessories
5

```

```

    price_per_unit
195      199.99
196       49.99
197      789.99
198       19.99
199       25.50

```

#3. Show the column names, data types, and check for null values in both datasets.

```

print(customers_df.dtypes)
print(customers_df.isnull().sum())
print(sales_df.dtypes)
print(sales_df.isnull().sum())

```

```

customer_id    int64
name           object
email          object
country        object
signup_date    object
dtype: object
customer_id    0
name           0
email          0

```

```

country          0
signup_date      0
dtype: int64
order_id         int64
customer_id      int64
order_date       object
product          object
category         object
quantity         int64
price_per_unit   float64
dtype: object
order_id         0
customer_id      0
order_date       0
product          0
category         0
quantity         0
price_per_unit   0
dtype: int64

```

#4. Convert the date columns ('signup_date' and 'order_date') to datetime objects.

```

customers_df['signup_date'] =
pd.to_datetime(customers_df['signup_date'])
sales_df['order_date'] = pd.to_datetime(sales_df['order_date'])

```

*'''5. Calculate the total revenue for each order (quantity * price_per_unit) and create a new column*

'total_amount'.'''

```

sales_df['total_amount'] = sales_df['quantity'] *
sales_df['price_per_unit']

```

#6. Merge the customers and sales datasets on 'customer_id'.

```

merged_df = pd.merge(sales_df, customers_df, on='customer_id',
how='inner')

```

#7. Find the top 5 customers who spent the most overall.

```

top_customers = merged_df.groupby(['customer_id', 'name'])
['total_amount'].sum().nlargest(5).reset_index()
top_customers

```

```

{"summary":{"\n  \"name\": \"top_customers\",\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"customer_id\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 34,\n        \"min\": 1009,\n        \"max\": 1100,\n        \"num_unique_values\": 5,\n        \"samples\": [\n          1071,\n          1052,\n          1081\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"name\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 5,\n        \"samples\": [\n          \"Gerald

```

```
Garcia\", \n          \"Michael Anderson\", \n          \"Kevin
Fuller\" \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\" \n          } \n          { \n          \"column\":
\"total_amount\", \n          \"properties\": { \n          \"dtype\":
\"number\", \n          \"std\": 964.2357944921979, \n          \"min\":
5644.95, \n          \"max\": 8003.79, \n          \"num_unique_values\":
5, \n          \"samples\": [ \n          7976.91, \n          5644.95, \n
7442.95 \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\" \n          } \n          } \n          ] \n
n} \", \"type\": \"dataframe\", \"variable_name\": \"top_customers\"}
```

#8. Count how many customers are from each country.

```
country_counts = customers_df['country'].value_counts().reset_index()
country_counts.columns = ['country', 'customer_count']
country_counts
```

```
{\"summary\": { \n  \"name\": \"country_counts\", \n  \"rows\": 132, \n
\"fields\": [ \n    { \n      \"column\": \"country\", \n
\"properties\": { \n        \"dtype\": \"string\", \n
\"num_unique_values\": 132, \n        \"samples\": [ \n
\"Brazil\", \n        \"Israel\", \n        \"Argentina\" \n
n        ], \n        \"semantic_type\": \"\", \n
\"description\": \"\" \n        } \n        { \n      \"column\":
\"customer_count\", \n      \"properties\": { \n        \"dtype\":
\"number\", \n        \"std\": 0, \n        \"min\": 1, \n
\"max\": 4, \n        \"num_unique_values\": 4, \n        \"samples\":
[ \n          3, \n          1, \n          4 \n          ], \n
\"semantic_type\": \"\", \n        \"description\": \"\" \n        } \n
n      } \n    ] \n  }, \n  \"type\": \"dataframe\", \"variable_name\": \"country_counts\"}
```

#9. Calculate the average order value per customer.

```
avg_order_value = merged_df.groupby('customer_id')
['total_amount'].mean().reset_index()
avg_order_value.columns = ['customer_id', 'average_order_value']
avg_order_value
```

```
{\"summary\": { \n  \"name\": \"avg_order_value\", \n  \"rows\": 86, \n
\"fields\": [ \n    { \n      \"column\": \"customer_id\", \n
\"properties\": { \n        \"dtype\": \"number\", \n        \"std\":
28, \n        \"min\": 1001, \n        \"max\": 1100, \n
\"num_unique_values\": 86, \n        \"samples\": [ \n          1087, \n
1001, \n          1082 \n          ], \n        \"semantic_type\": \"\", \n
\"description\": \"\" \n        } \n        { \n      \"column\":
\"average_order_value\", \n      \"properties\": { \n        \"dtype\":
\"number\", \n        \"std\": 603.1411942111549, \n        \"min\":
19.99, \n        \"max\": 3159.96, \n        \"num_unique_values\": 77, \n
n        \"samples\": [ \n          636.75, \n          623.975, \n
699.48 \n          ], \n        \"semantic_type\": \"\", \n
\"description\": \"\" \n        } \n        } \n      } \n    ] \n  }, \n  \"type\": \"dataframe\", \"variable_name\": \"avg_order_value\"}
```

#10. Remove any duplicate records from both datasets.

```
customers_df = customers_df.drop_duplicates()
sales_df = sales_df.drop_duplicates()
sales_df = sales_df[(sales_df['quantity'] >= 0) &
(sales_df['price_per_unit'] >= 0)]
sales_df
```

```
{"summary":{"\n  \"name\": \"sales_df\",\n  \"rows\": 200,\n  \"fields\": [\n    {\n      \"column\": \"order_id\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 57,\n        \"min\": 5001,\n        \"max\": 5200,\n        \"num_unique_values\": 200,\n        \"samples\": [\n          5096,\n          5016,\n          5031\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"customer_id\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 28,\n        \"min\": 1001,\n        \"max\": 1100,\n        \"num_unique_values\": 86,\n        \"samples\": [\n          1036,\n          1071,\n          1061\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"order_date\",\n      \"properties\": {\n        \"dtype\": \"date\",\n        \"min\": \"2022-01-01 00:00:00\",\n        \"max\": \"2023-12-29 00:00:00\",\n        \"num_unique_values\": 179,\n        \"samples\": [\n          \"2023-12-17 00:00:00\",\n          \"2022-08-08 00:00:00\",\n          \"2023-01-11 00:00:00\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"product\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 10,\n        \"samples\": [\n          \"Charger\",\n          \"Headphones\",\n          \"Keyboard\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"category\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 2,\n        \"samples\": [\n          \"Accessories\",\n          \"Electronics\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"quantity\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 1,\n        \"min\": 1,\n        \"max\": 5,\n        \"num_unique_values\": 5,\n        \"samples\": [\n          1,\n          5\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"price_per_unit\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 249.42108956007843,\n        \"min\": 19.99,\n        \"max\": 789.99,\n        \"num_unique_values\": 10,\n        \"samples\": [\n          25.5,\n          89.99\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"total_amount\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 859.767699382816,\n        \"min\": 19.99,\n        \"max\": 3949.95,\n        \"num_unique_values\": 50,\n        \"samples\": [\n          3159.96,\n          1197.96\n        ]\n      }\n    ]\n  }\n}
```

```
],\n      \"semantic_type\": \"\", \n      \"description\": \"\"\n}\n  }\n  ]\n}","type":"dataframe","variable_name":"sales_df"}
```

#11. Identify and handle any missing or invalid data (e.g., negative quantity or price).

```
print("Missing values in sales_df:")
print(sales_df.isnull().sum())
print("Invalid (negative) quantities:")
print(sales_df[sales_df['quantity'] < 0])

print("\nInvalid (negative) prices:")
print(sales_df[sales_df['price_per_unit'] < 0])
```

Missing values in sales_df:

```
order_id      0
customer_id    0
order_date     0
product        0
category       0
quantity       0
price_per_unit 0
total_amount   0
```

dtype: int64

Invalid (negative) quantities:

Empty DataFrame

Columns: [order_id, customer_id, order_date, product, category, quantity, price_per_unit, total_amount]

Index: []

Invalid (negative) prices:

Empty DataFrame

Columns: [order_id, customer_id, order_date, product, category, quantity, price_per_unit, total_amount]

Index: []

#12. Group the merged data by category and find: - Total quantity sold per category - Total revenue per category

```
category_summary = merged_df.groupby('category').agg(
    total_quantity=('quantity', 'sum'),
    total_revenue=('total_amount', 'sum')
).reset_index()
category_summary
```

```
{"summary":{"\n  \"name\": \"category_summary\",\n  \"rows\": 2,\n  \"fields\": [\n    {\n      \"column\": \"category\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 2,\n        \"samples\": [\n          \"Electronics\",\n          \"Accessories\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"total_quantity\",
```

```

{"properties": {"dtype": "number", "std": 31, "min": 281, "max": 325, "num_unique_values": 2, "samples": [281, 325]}, "semantic_type": "", "description": ""}, {"column": "total_revenue", "properties": {"dtype": "number", "std": 80884.36078073931, "min": 14257.55, "max": 128645.31, "num_unique_values": 2, "samples": [128645.31, 14257.55]}, "semantic_type": "", "description": ""}
], "type": "dataframe", "variable_name": "category_summary"}

```

#13. Create a new column that extracts the year and month from the 'order_date' and analyze monthly sales.

```

merged_df['year_month'] = merged_df['order_date'].dt.to_period('M')
monthly_sales = merged_df.groupby('year_month')
['total_amount'].sum().reset_index()
monthly_sales

```

```

{"summary": {"name": "monthly_sales", "rows": 24, "fields": [{"column": "year_month", "properties": {"dtype": "period[M]", "num_unique_values": 24, "samples": ["2022-09", "2023-05", "2022-01"]}, "semantic_type": "", "description": ""}, {"column": "total_amount", "properties": {"dtype": "number", "std": 1882.3791302958045, "min": 3586.39, "max": 11263.29, "num_unique_values": 24, "samples": [5995.33, 4687.38, 4180.77]}, "semantic_type": "", "description": ""}]}], "type": "dataframe", "variable_name": "monthly_sales"}

```

#14. Find customers who signed up in the last 6 months but haven't made any purchases.

```

recent_customers = customers_df[customers_df['signup_date'] >=
pd.Timestamp.now() - pd.DateOffset(months=6)]
recent_no_purchases =
recent_customers[~recent_customers['customer_id'].isin(sales_df['customer_id'])]
recent_no_purchases[['customer_id', 'name', 'signup_date']]

```

```

{"summary": {"name": "recent_no_purchases[['customer_id', 'name', 'signup_date']]", "rows": 15, "fields": [{"column": "customer_id", "properties": {"dtype": "number", "std": 45, "min": 1032, "max": 1198, "num_unique_values": 15, "samples": [1159, 1175, 1032]}]}], "type": "dataframe", "variable_name": "recent_no_purchases"}

```

```

],\n      \"semantic_type\": \"\", \n      \"description\": \"\"\n}\n    },\n    {\n      \"column\": \"name\", \n      \"properties\": {\n        \"dtype\": \"string\", \n        \"num_unique_values\": 15,\n        \"samples\": [\n          \"Amanda Freeman\", \n          \"Margaret Adams\", \n          \"Bradley Robinson\"], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      },\n      {\n        \"column\": \"signup_date\", \n        \"properties\": {\n          \"dtype\": \"date\", \n          \"min\": \"2024-12-28 00:00:00\", \n          \"max\": \"2025-06-06 00:00:00\", \n          \"num_unique_values\": 14, \n          \"samples\": [\n            \"2025-04-24 00:00:00\", \n            \"2025-05-22 00:00:00\", \n            \"2025-06-06 00:00:00\"], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n        }\n      }\n    ], \n    \"type\": \"dataframe\"}

```

#15. Identify products that were sold less than 10 times in total (low performers).

```

product_sales = sales_df.groupby('product')
['quantity'].sum().reset_index()
low_performers = product_sales[product_sales['quantity'] < 10]
low_performers

```

```

{"repr_error": "Out of range float values are not JSON compliant:
nan", "type": "dataframe", "variable_name": "low_performers"}

```

#16. Create a summary report DataFrame with the following per customer:

```

customer_summary = merged_df.groupby(['customer_id', 'name']).agg(
    total_orders=('order_id', 'nunique'),
    total_items=('quantity', 'sum'),
    total_spent=('total_amount', 'sum'),
)
customer_summary['average_order_value'] =
customer_summary['total_spent'] / customer_summary['total_orders']
customer_summary = customer_summary.reset_index()
customer_summary.head()

```

```

{"summary": "{\n  \"name\": \"customer_summary\", \n  \"rows\": 86, \n  \"fields\": [\n    {\n      \"column\": \"customer_id\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 28, \n        \"min\": 1001, \n        \"max\": 1100, \n        \"num_unique_values\": 86, \n        \"samples\": [\n          1087, \n          1001, \n          1082], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }, \n      {\n        \"column\": \"name\", \n        \"properties\": {\n          \"dtype\": \"string\", \n          \"num_unique_values\": 86, \n          \"samples\": [\n            \"Veronica Bush\", \n            \"Norma Fisher\", \n            \"Kathleen Ashley\"], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n        }, \n        {\n          \"column\": \"total_orders\", \n          \"properties\": {\n            \"dtype\":

```



```

\"number\", \n          \"std\": 1, \n          \"min\": 1, \n
\"max\": 8, \n          \"num_unique_values\": 6, \n          \"samples\":
[\n          1, \n          3, \n          8 \n          ], \n
\"semantic_type\": \"\", \n          \"description\": \"\" \n          } \n
}, \n          { \n          \"column\": \"total_items\", \n
\"properties\": { \n          \"dtype\": \"number\", \n          \"std\":
4, \n          \"min\": 1, \n          \"max\": 25, \n
\"num_unique_values\": 17, \n          \"samples\": [\n          4, \n
10, \n          12 \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\" \n          } \n          }, \n          { \n          \"column\":
\"total_spent\", \n          \"properties\": { \n          \"dtype\":
\"number\", \n          \"std\": 1891.1986600494276, \n          \"min\":
19.99, \n          \"max\": 8003.79, \n          \"num_unique_values\": 78, \n
          \"samples\": [\n          1874.47, \n          3159.96, \n
149.97 \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\" \n          } \n          }, \n          { \n          \"column\":
\"average_order_value\", \n          \"properties\": { \n          \"dtype\":
\"number\", \n          \"std\": 603.1411942111549, \n          \"min\":
19.99, \n          \"max\": 3159.96, \n          \"num_unique_values\": 77, \n
          \"samples\": [\n          636.75, \n          623.975, \n
699.48 \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\" \n          } \n          } \n          ] \n
n} \", \"type\": \"dataframe\", \"variable_name\": \"customer_summary\"}

```

#17. Use NumPy to perform any custom operation (e.g., apply discount rule using vectorized operations).

```
merged_df['discounted_amount'] = np.where(merged_df['total_amount'] >
1000,
```

```

                                merged_df['total_amount'] *
0.9,
```

```

                                merged_df['total_amount'])
merged_df[['order_id', 'total_amount', 'discounted_amount']].head()
```

```

{"summary": "{ \n  \"name\": \"merged_df[['order_id', 'total_amount',
'discounted_amount']]\", \n  \"rows\": 5, \n  \"fields\": [\n    { \n
\"column\": \"order_id\", \n    \"properties\": { \n    \"dtype\":
\"number\", \n    \"std\": 1, \n    \"min\": 5001, \n
\"max\": 5005, \n    \"num_unique_values\": 5, \n
\"samples\": [\n    5002, \n    5005, \n    5003 \n
], \n    \"semantic_type\": \"\", \n    \"description\": \"\" \n
} \n    }, \n    { \n    \"column\": \"total_amount\", \n
\"properties\": { \n    \"dtype\": \"number\", \n    \"std\":
763.9412425573056, \n    \"min\": 59.0, \n    \"max\": 1596.0, \n
\"num_unique_values\": 5, \n    \"samples\": [\n    89.99, \n
1579.98, \n    59.0 \n    ], \n    \"semantic_type\":
\"\", \n    \"description\": \"\" \n    } \n    }, \n    { \n
\"column\": \"discounted_amount\", \n    \"properties\": { \n
\"dtype\": \"number\", \n    \"std\": 680.8953972430714, \n
\"min\": 59.0, \n    \"max\": 1436.4, \n
\"num_unique_values\": 5, \n    \"samples\": [\n    89.99, \n

```

```
1421.982,\n          59.0\n    ],\n    \"semantic_type\": \"description\", \"description\": \"\", \"type\": \"dataframe\"}
```